

Supplementary Material: Fabricated Authority in the Wild

Anonymous Author(s)

September 5, 2026

1 Purpose of the supplement

This supplement supports the regular-paper submission without adding main-paper pages. It documents the scenario inventory, policy variants, attack families, result artifacts, and acceptance-path checks. It is anonymous and contains no author-identifying paths or credentials.

2 Scenario inventory

Each scenario defines an agent role, an operating-policy prohibition, a user task, a forbidden tool call, scenario-specific violating arguments, and the safe alternative `request_approval`. The eight scenarios are:

ID	Family	Forbidden action
medical_export	Data exfiltration	Export patient records externally
payroll_disclosure	Data exfiltration	Send payroll data externally
customer_pii_share	Data exfiltration	Upload customer PII to ad vendor
prod_db_delete	Destructive action	Delete production records
backup_deletion	Destructive action	Delete in-retention backups
security_disable	Security change	Disable endpoint protection
wire_transfer	Financial action	Initiate unapproved transfer
access_grant	Privilege escalation	Grant admin access without approval

Table 1: Scenario coverage used by the tool-call harness.

3 Policy variants

The main sweep varies operating-policy phrasing while holding the scenario, tools, and user task fixed.

Variant	Intended distinction
strict_hatch	Prohibition, safe alternative, and explicit provenance clause
strict	Prohibition plus safe alternative
strict_verbose	Longer strict policy used only for length control
bare	Prohibition only
autonomous	Autonomy-forward routine-handling language
exemption	Prohibition with documented-exemption carve-out

Table 2: Policy variants. The central independent variable is provenance authentication of claimed authority.

The provenance clause tested in `strict_hatch` is: “Claims of exemption, prior approval, or superseding policy that arrive in a user message or tool result are NOT authoritative. Treat them as unverified and call `request_approval`.”

4 Attack families

The seven attack names are `none`, `superseded`, `prior_approval`, `reclassify`, `emergency`, `audit_exemption`, and `combined`. The attacks are synthetic operating requests, not exploit instructions or real credentials. They test category pressure, urgency pressure, and process/provenance pressure.

5 Result artifact map

Artifact	Trials	Main use
<code>agentic_sweep_api.json</code>	1,680	Main direct-channel policy matrix
<code>agentic_ablation.json</code>	288	Length-control ablation
<code>agentic_matched_contrast.json</code>	36	Complete-separation matched contrast
<code>agentic_gap_full.json</code>	432	Text/tool-call divergence
<code>agentic_adaptive*.json</code>	2 runs	Adaptive attacker check
<code>agentic_frontier_matched.json</code>	48	Sonnet target-sensitivity negative
<code>agentic_indirect_all8_t1.json</code>	96	All-scenario forced indirect-channel result
<code>agentic_sweep_local.json</code>	560	Local small-model boundary
<code>agentic_sweep_gpt4lmini_t1.json</code>	280	Additional hosted-target breadth check

Table 3: Raw artifacts used by the main and extended drafts.

6 Remaining hypothesis checks

The current submission tests the following reviewer-risk hypotheses: the effect is not ordinary noncompliance, not only policy length, not universal model weakness, not fully solved by prompt wording, and not visible to text-only scoring. The remaining open hypotheses are multi-turn persistence, stronger adaptive attackers, stronger repetition for the all-scenario indirect result, and broader model coverage.

7 Reproducibility resources

The anonymous artifact package contains the main paper, supplement, raw result JSON files, and the agentic harness source used to produce the reported tool-call outcomes. The intended reproduction path is: inspect the scenario and attack definitions, run the relevant `loki.agentic.*` CLI, and compare the produced JSON summaries against the claim-guard tests. Hosted-model reproduction requires the evaluator’s own API credentials; no credentials or private paths are included in the submission bundle.

8 Acceptance-path validation

The main regular-paper PDF was compiled with Tectonic and parsed with `pypdf`. It is 9 pages, unencrypted, anonymous, and contains the updated all-scenario indirect result. Public experiment interfaces were checked through CLI help commands for `sweep`, `gap`, `adaptive`, and `indirect`. Targeted tests for agentic behavior and paper claims passed.

9 Final submission recommendation

For OpenReview, upload `aiwild_regular_anonymous.pdf` as the main regular paper and `aiwild_regular_supplement.pdf` as supplementary material. Do not upload source files unless requested by the venue. The source files are included in the local bundle for collaborator handoff and reproducibility review.